

```
In [1]: pip install gensim
```

```
Requirement already satisfied: gensim in /usr/local/lib/python3.12/dist-packages (4.3.3)  
Requirement already satisfied: numpy<2.0,>=1.18.5 in /usr/local/lib/python3.12/dist-packages (from gensim) (1.26.4)  
Requirement already satisfied: scipy<1.14.0,>=1.7.0 in /usr/local/lib/python3.12/dist-packages (from gensim) (1.13.1)  
Requirement already satisfied: smart-open>=1.8.1 in /usr/local/lib/python3.12/dist-packages (from gensim) (7.3.1)  
Requirement already satisfied: wrapt in /usr/local/lib/python3.12/dist-packages (from smart-open>=1.8.1->gensim) (1.17.3)
```

```
In [22]: import pandas as pd  
from datasets import load_dataset  
from imblearn.under_sampling import RandomUnderSampler  
import nltk  
from nltk.corpus import stopwords  
import re  
from sklearn.model_selection import train_test_split  
import numpy as np  
from sklearn.feature_extraction.text import TfidfVectorizer  
from sklearn.linear_model import LogisticRegression  
from sklearn.svm import LinearSVC  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score  
from dataclasses import dataclass  
from typing import List, Tuple  
from typing import List  
import tensorflow as tf  
from tensorflow.keras.preprocessing.text import Tokenizer  
from tensorflow.keras.preprocessing.sequence import pad_sequences  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import (  
    Embedding,  
    LSTM,  
    Dense,  
    Conv1D,  
    MaxPooling1D,  
    Dropout,  
    GlobalMaxPooling1D,  
)  
from transformers import AutoTokenizer, TFAutoModelForSequenceClassification  
  
import gensim.downloader as api
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

Classes

```
In [3]: @dataclass
class ModelResult:
    """Simple container for model evaluation results."""
    name: str
    accuracy: float
    precision: float
    recall: float
    f1: float
```

Data Selection

```
In [4]: seed = 42
sample_size = 5000
# Load dataset using huggingface datasets library
dataset = load_dataset('ucirvine/sms_spam', split='train')
df = pd.DataFrame(dataset)
# Rename column for convenience
df = df.rename(columns={'sms': 'text'})
# Sample if larger than allowed
if len(df) > sample_size:
    df = df.sample(n=sample_size, random_state=seed).reset_index(drop=True)
df.head()
```

/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:

The secret `HF_TOKEN` does not exist in your Colab secrets.

To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>), set it as secret in your Google Colab and restart your session.

You will be able to reuse this secret in all of your notebooks.

Please note that authentication is recommended but still optional to access public models or datasets.

warnings.warn(

Out[4]:

	text	label
0	You still coming tonight?\n	0
1	"HEY BABE! FAR 2 SPUN-OUT 2 SPK AT DA MO... DE...	0
2	Ya even those cookies have jelly on them\n	0
3	Sorry i've not gone to that place. I'll do so ...	0
4	When are you going to ride your bike?\n	0

Preprocessing

```
In [5]: # Ensure stopwords are downloaded
        nltk.download('stopwords')

        def clean_text(text: str, stop_words: set) -> str:
            text = text.lower()
            text = re.sub(r'http\S+', ' ', text) # remove URLs
            text = re.sub(r'\S+@\S+', ' ', text) # remove emails
            text = re.sub(r'\d+', ' ', text) # remove numbers
            text = re.sub(r'[^\w-zA-Z]', ' ', text) # remove punctuation and special chars
            tokens = text.split()
            tokens = [t for t in tokens if t not in stop_words]
            return ' '.join(tokens)

        english_stopwords = set(stopwords.words('english'))
        df['clean_text'] = df['text'].apply(lambda x: clean_text(x, english_stopwords))
        df.head()
```

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!

Out[5]:

	text	label	clean_text
0	You still coming tonight?\n	0	still coming tonight
1	"HEY BABE! FAR 2 SPUN-OUT 2 SPK AT DA MO... DE...	0	hey babe far spun spk da mo dead da wrld sleep...
2	Ya even those cookies have jelly on them\n	0	ya even cookies jelly
3	Sorry i've not gone to that place. I'll do so ...	0	sorry gone place tomorrow really sorry
4	When are you going to ride your bike?\n	0	going ride bike

Balance the dataset

```
In [6]: def balance_dataset(texts, labels, random_state = 42):
        rus = RandomUnderSampler(random_state=random_state)
        X = pd.DataFrame({'text': texts})
        y = pd.Series(labels)
        X_resampled, y_resampled = rus.fit_resample(X, y)
        return X_resampled['text'].tolist(), y_resampled.tolist()

balanced_texts, balanced_labels = balance_dataset(df['clean_text'].tolist(), df['label'].tolist())
```

Split into train and test sets

```
In [7]: X_train_texts, X_test_texts, y_train, y_test = train_test_split(
        balanced_texts,
        balanced_labels,
        test_size=0.2,
        random_state=42,
        stratify=balanced_labels,
    )
```

Traditional ML models

```

In [8]: def vectorize_texts(train_texts: List[str], test_texts: List[str]) -> Tuple[np.ndarray, np.ndarray, TfidfVectorizer]:
        vectorizer = TfidfVectorizer(
            max_features=5000, ngram_range=(1, 2), sublinear_tf=True, stop_words='english'
        )
        X_train = vectorizer.fit_transform(train_texts)
        X_test = vectorizer.transform(test_texts)
        return X_train, X_test, vectorizer

def train_ml_models(
    X_train: np.ndarray,
    X_test: np.ndarray,
    y_train: List[int],
    y_test: List[int],
) -> List[ModelResult]:
    results = []
    # Logistic Regression
    log_reg = LogisticRegression(max_iter=1000)
    log_reg.fit(X_train, y_train)
    y_pred = log_reg.predict(X_test)
    results.append(
        ModelResult(
            name='Logistic Regression',
            accuracy=accuracy_score(y_test, y_pred),
            precision=precision_score(y_test, y_pred, zero_division=0),
            recall=recall_score(y_test, y_pred, zero_division=0),
            f1=f1_score(y_test, y_pred, zero_division=0),
        )
    )

    # Linear SVM
    svm_model = LinearSVC()
    svm_model.fit(X_train, y_train)
    y_pred = svm_model.predict(X_test)
    results.append(
        ModelResult(
            name='Linear SVM',
            accuracy=accuracy_score(y_test, y_pred),
            precision=precision_score(y_test, y_pred, zero_division=0),
            recall=recall_score(y_test, y_pred, zero_division=0),
            f1=f1_score(y_test, y_pred, zero_division=0),
        )
    )

```

```
    )
)

# Random Forest
rf_model = RandomForestClassifier(n_estimators=200, random_state=42)
rf_model.fit(X_train, y_train)
y_pred = rf_model.predict(X_test)
results.append(
    ModelResult(
        name='Random Forest',
        accuracy=accuracy_score(y_test, y_pred),
        precision=precision_score(y_test, y_pred, zero_division=0),
        recall=recall_score(y_test, y_pred, zero_division=0),
        f1=f1_score(y_test, y_pred, zero_division=0),
    )
)
return results
```

```
In [9]: X_train_tfidf, X_test_tfidf, _ = vectorize_texts(X_train_texts, X_test_texts)
ml_results = train_ml_models(X_train_tfidf, X_test_tfidf, y_train, y_test)
```

```
In [10]: ml_results
```

```
Out[10]: [ModelResult(name='Logistic Regression', accuracy=0.9547169811320755, precision=0.9615384615384616, recall=0.946969696969697, f1=0.9541984732824428),
ModelResult(name='Linear SVM', accuracy=0.9547169811320755, precision=0.9615384615384616, recall=0.946969696969697, f1=0.9541984732824428),
ModelResult(name='Random Forest', accuracy=0.9320754716981132, precision=0.9523809523809523, recall=0.9090909090909091, f1=0.9302325581395349)]
```

Deep learning models

```
In [14]: def build_lstm_model(
vocab_size: int,
embedding_dim: int = 128,
input_length: int = 100,
) -> tf.keras.Model:
    """
    Build and compile a simple LSTM model for text classification.

    Parameters
    -----
```

```

vocab_size : int
    Size of the vocabulary (number of unique tokens).
embedding_dim : int
    Dimensionality of the embedding space.
input_length : int
    Length of the input sequences (after padding/truncation).

Returns
-----
model : tf.keras.Model
    The compiled LSTM model.
"""
model = Sequential([
    Embedding(input_dim=vocab_size, output_dim=embedding_dim, input_length=input_length),
    LSTM(128, return_sequences=False),
    Dropout(0.5),
    Dense(64, activation='relu'),
    Dropout(0.5),
    Dense(1, activation='sigmoid'),
])
model.compile(
    loss='binary_crossentropy',
    optimizer='adam',
    metrics=['accuracy', tf.keras.metrics.Precision(), tf.keras.metrics.Recall()],
)
return model

def build_cnn_model(
    vocab_size: int,
    embedding_dim: int = 128,
    input_length: int = 100,
) -> tf.keras.Model:
    """
    Build and compile a 1D CNN model for text classification.

    Parameters
    -----
    vocab_size : int
        Size of the vocabulary.
    embedding_dim : int
        Dimensionality of the embedding space.

```

```

input_length : int
    Length of the input sequences.

Returns
-----
model : tf.keras.Model
    The compiled CNN model.
"""
model = Sequential([
    Embedding(input_dim=vocab_size, output_dim=embedding_dim, input_length=input_length),
    Conv1D(filters=128, kernel_size=5, activation='relu'),
    MaxPooling1D(pool_size=2),
    Dropout(0.5),
    Conv1D(filters=64, kernel_size=3, activation='relu'),
    GlobalMaxPooling1D(),
    Dropout(0.5),
    Dense(64, activation='relu'),
    Dropout(0.5),
    Dense(1, activation='sigmoid'),
])
model.compile(
    loss='binary_crossentropy',
    optimizer='adam',
    metrics=['accuracy', tf.keras.metrics.Precision(), tf.keras.metrics.Recall()],
)
return model

def train_deep_models(
    train_texts: List[str],
    test_texts: List[str],
    y_train: List[int],
    y_test: List[int],
    max_num_words: int = 10000,
    max_sequence_length: int = 100,
    batch_size: int = 32,
    lstm_epochs: int = 5,
    cnn_epochs: int = 5,
    bert_epochs: int = 2,
    bert_model_name: str = 'bert-base-uncased',
) -> List[ModelResult]:
    """

```


Train deep learning models (LSTM, CNN, BERT) and evaluate them.

Parameters

train_texts : List[str]

Cleaned training texts.

test_texts : List[str]

Cleaned testing texts.

y_train : List[int]

Training labels.

y_test : List[int]

Testing labels.

max_num_words : int

Maximum number of words to keep in the tokenizer vocabulary.

max_sequence_length : int

Maximum length of sequences (padding/truncation).

batch_size : int

Batch size for training deep models.

lstm_epochs : int

Epochs for LSTM model.

cnn_epochs : int

Epochs for CNN model.

bert_epochs : int

Epochs for BERT model.

bert_model_name : str

Hugging Face model name (e.g. 'bert-base-uncased').

Returns

results : List[ModelResult]

Evaluation results for each model.

"""

results = []

Convert texts to sequences for LSTM and CNN

tokenizer = Tokenizer(num_words=max_num_words, oov_token='<OOV>')

tokenizer.fit_on_texts(train_texts)

X_train_seq = tokenizer.texts_to_sequences(train_texts)

X_test_seq = tokenizer.texts_to_sequences(test_texts)

X_train_pad = pad_sequences(X_train_seq, maxlen=max_sequence_length, padding='post', truncating='post')

X_test_pad = pad_sequences(X_test_seq, maxlen=max_sequence_length, padding='post', truncating='post')

y_train_array = np.array(y_train)

```

y_test_array = np.array(y_test)
vocab_size = min(max_num_words, len(tokenizer.word_index) + 1)

# -----
# Train LSTM (random embeddings)
# -----
lstm_model = build_lstm_model(vocab_size=vocab_size, input_length=max_sequence_length)
lstm_model.fit(
    X_train_pad,
    y_train_array,
    epochs=lstm_epochs,
    batch_size=batch_size,
    validation_split=0.1,
    verbose=2,
)
lstm_metrics = lstm_model.evaluate(X_test_pad, y_test_array, verbose=0)
lstm_accuracy = lstm_metrics[1]
lstm_precision = lstm_metrics[2]
lstm_recall = lstm_metrics[3]
lstm_f1 = 2 * (lstm_precision * lstm_recall) / (lstm_precision + lstm_recall + 1e-8)
results.append(
    ModelResult(
        name='LSTM',
        accuracy=lstm_accuracy,
        precision=lstm_precision,
        recall=lstm_recall,
        f1=lstm_f1,
    )
)

# -----
# Train LSTM with GloVe embeddings
# -----
# Load pre-trained GloVe embeddings (100-dimensional). This step can be
# time-consuming; consider caching the embedding model locally.
try:
    glove_model = api.load('glove-wiki-gigaword-100')
except Exception as e:
    print(f"Warning: failed to load GloVe embeddings: {e}. Skipping pretrained LSTM model.")
    glove_model = None

if glove_model is not None:

```

```

embedding_dim = 100
# Initialize embedding matrix with zeros
embedding_matrix = np.zeros((vocab_size, embedding_dim))
for word, i in tokenizer.word_index.items():
    if i >= vocab_size:
        continue
    if word in glove_model:
        embedding_matrix[i] = glove_model[word]
# Build LSTM model with pre-trained embeddings
lstm_pretrained = Sequential([
    Embedding(
        input_dim=vocab_size,
        output_dim=embedding_dim,
        weights=[embedding_matrix],
        input_length=max_sequence_length,
        trainable=False,
    ),
    LSTM(128, return_sequences=False),
    Dropout(0.5),
    Dense(64, activation='relu'),
    Dropout(0.5),
    Dense(1, activation='sigmoid'),
])
lstm_pretrained.compile(
    loss='binary_crossentropy',
    optimizer='adam',
    metrics=['accuracy', tf.keras.metrics.Precision(), tf.keras.metrics.Recall()],
)
lstm_pretrained.fit(
    X_train_pad,
    y_train_array,
    epochs=lstm_epochs,
    batch_size=batch_size,
    validation_split=0.1,
    verbose=2,
)
glove_metrics = lstm_pretrained.evaluate(X_test_pad, y_test_array, verbose=0)
glove_accuracy = glove_metrics[1]
glove_precision = glove_metrics[2]
glove_recall = glove_metrics[3]
glove_f1 = 2 * (glove_precision * glove_recall) / (glove_precision + glove_recall + 1e-8)
results.append(

```

```

        ModelResult(
            name='LSTM + GloVe',
            accuracy=glove_accuracy,
            precision=glove_precision,
            recall=glove_recall,
            f1=glove_f1,
        )
    )

# Train CNN
cnn_model = build_cnn_model(vocab_size=vocab_size, input_length=max_sequence_length)
cnn_model.fit(
    X_train_pad,
    y_train_array,
    epochs=cnn_epochs,
    batch_size=batch_size,
    validation_split=0.1,
    verbose=2,
)
cnn_metrics = cnn_model.evaluate(X_test_pad, y_test_array, verbose=0)
cnn_accuracy = cnn_metrics[1]
cnn_precision = cnn_metrics[2]
cnn_recall = cnn_metrics[3]
cnn_f1 = 2 * (cnn_precision * cnn_recall) / (cnn_precision + cnn_recall + 1e-8)
results.append(
    ModelResult(
        name='CNN',
        accuracy=cnn_accuracy,
        precision=cnn_precision,
        recall=cnn_recall,
        f1=cnn_f1,
    )
)

# -----
# Train a custom Transformer model using Keras
# -----
# Build a transformer-inspired network with multi-head attention. This avoids
# external dependencies on Hugging Face models and satisfies the
# transformer-based requirement.
from tensorflow.keras.layers import MultiHeadAttention, LayerNormalization, GlobalAveragePooling1D, Input
from tensorflow.keras.models import Model

```

```

inputs = Input(shape=(max_sequence_length,))
x_emb = Embedding(input_dim=vocab_size, output_dim=128)(inputs)
# Multi-head self-attention
attn_output = MultiHeadAttention(num_heads=4, key_dim=32)(x_emb, x_emb)
# Add & norm
x = LayerNormalization(epsilon=1e-6)(attn_output + x_emb)
# Global average pooling
x = GlobalAveragePooling1D()(x)
x = Dropout(0.5)(x)
x = Dense(64, activation='relu')(x)
x = Dropout(0.5)(x)
outputs = Dense(1, activation='sigmoid')(x)
transformer_model = Model(inputs=inputs, outputs=outputs)
transformer_model.compile(
    loss='binary_crossentropy',
    optimizer='adam',
    metrics=['accuracy', tf.keras.metrics.Precision(), tf.keras.metrics.Recall()],
)
transformer_model.fit(
    X_train_pad,
    y_train_array,
    epochs=cnn_epochs, # reuse cnn_epochs for transformer
    batch_size=batch_size,
    validation_split=0.1,
    verbose=2,
)
transformer_metrics = transformer_model.evaluate(X_test_pad, y_test_array, verbose=0)
transformer_accuracy = transformer_metrics[1]
transformer_precision = transformer_metrics[2]
transformer_recall = transformer_metrics[3]
transformer_f1 = 2 * (transformer_precision * transformer_recall) / (transformer_precision + transformer_recall + 1e-8)
results.append(
    ModelResult(
        name='Transformer',
        accuracy=transformer_accuracy,
        precision=transformer_precision,
        recall=transformer_recall,
        f1=transformer_f1,
    )
)
return results

```

```
In [15]: deep_results = train_deep_models(  
    train_texts=X_train_texts,  
    test_texts=X_test_texts,  
    y_train=y_train,  
    y_test=y_test,  
    max_num_words=10000,  
    max_sequence_length=100,  
    batch_size=32,  
    lstm_epochs=5,  
    cnn_epochs=5,  
    bert_epochs=2,  
    bert_model_name='bert-base-uncased',  
)
```

Epoch 1/5
30/30 - 11s - 369ms/step - accuracy: 0.5047 - loss: 0.6980 - precision_3: 0.5158 - recall_3: 0.3727 - val_accuracy: 0.4434 - val_loss: 0.7007 - val_precision_3: 0.4434 - val_recall_3: 1.0000

Epoch 2/5
30/30 - 10s - 328ms/step - accuracy: 0.5005 - loss: 0.6954 - precision_3: 0.5047 - recall_3: 0.7723 - val_accuracy: 0.5566 - val_loss: 0.6919 - val_precision_3: 0.0000e+00 - val_recall_3: 0.0000e+00

Epoch 3/5
30/30 - 5s - 165ms/step - accuracy: 0.4858 - loss: 0.6938 - precision_3: 0.4927 - recall_3: 0.4865 - val_accuracy: 0.5566 - val_loss: 0.6926 - val_precision_3: 0.0000e+00 - val_recall_3: 0.0000e+00

Epoch 4/5
30/30 - 6s - 196ms/step - accuracy: 0.5016 - loss: 0.6956 - precision_3: 0.5091 - recall_3: 0.4617 - val_accuracy: 0.4434 - val_loss: 0.6957 - val_precision_3: 0.4434 - val_recall_3: 1.0000

Epoch 5/5
30/30 - 6s - 187ms/step - accuracy: 0.5026 - loss: 0.6939 - precision_3: 0.5066 - recall_3: 0.7205 - val_accuracy: 0.4434 - val_loss: 0.6945 - val_precision_3: 0.4434 - val_recall_3: 1.0000

Epoch 1/5
30/30 - 10s - 343ms/step - accuracy: 0.4890 - loss: 0.6941 - precision_4: 0.4976 - recall_4: 0.8468 - val_accuracy: 0.4434 - val_loss: 0.6937 - val_precision_4: 0.4434 - val_recall_4: 1.0000

Epoch 2/5
30/30 - 4s - 136ms/step - accuracy: 0.4711 - loss: 0.6937 - precision_4: 0.4862 - recall_4: 0.7660 - val_accuracy: 0.4434 - val_loss: 0.6941 - val_precision_4: 0.4434 - val_recall_4: 1.0000

Epoch 3/5
30/30 - 5s - 181ms/step - accuracy: 0.4921 - loss: 0.6939 - precision_4: 0.4993 - recall_4: 0.7702 - val_accuracy: 0.4434 - val_loss: 0.6938 - val_precision_4: 0.4434 - val_recall_4: 1.0000

Epoch 4/5
30/30 - 4s - 132ms/step - accuracy: 0.5068 - loss: 0.6936 - precision_4: 0.5071 - recall_4: 0.9627 - val_accuracy: 0.4434 - val_loss: 0.6955 - val_precision_4: 0.4434 - val_recall_4: 1.0000

Epoch 5/5
30/30 - 4s - 137ms/step - accuracy: 0.5026 - loss: 0.6931 - precision_4: 0.5049 - recall_4: 0.9669 - val_accuracy: 0.4434 - val_loss: 0.6949 - val_precision_4: 0.4434 - val_recall_4: 1.0000

Epoch 1/5
30/30 - 6s - 208ms/step - accuracy: 0.5205 - loss: 0.6924 - precision_5: 0.5203 - recall_5: 0.6894 - val_accuracy: 0.7075 - val_loss: 0.6681 - val_precision_5: 0.6143 - val_recall_5: 0.9149

Epoch 2/5
30/30 - 2s - 58ms/step - accuracy: 0.7261 - loss: 0.5889 - precision_5: 0.6672 - recall_5: 0.9172 - val_accuracy: 0.8019 - val_loss: 0.4726 - val_precision_5: 0.7167 - val_recall_5: 0.9149

Epoch 3/5
30/30 - 2s - 58ms/step - accuracy: 0.9486 - loss: 0.2222 - precision_5: 0.9617 - recall_5: 0.9358 - val_accuracy: 0.9151 - val_loss: 0.3655 - val_precision_5: 0.9130 - val_recall_5: 0.8936

Epoch 4/5
30/30 - 2s - 58ms/step - accuracy: 0.9738 - loss: 0.1114 - precision_5: 0.9872 - recall_5: 0.9607 - val_accuracy: 0.8962 - val_loss: 0.4166 - val_precision_5: 0.8913 - val_recall_5: 0.8723

```
Epoch 5/5
30/30 - 3s - 89ms/step - accuracy: 0.9822 - loss: 0.0779 - precision_5: 0.9936 - recall_5: 0.9710 - val_accuracy: 0.8491 - val_loss: 0.6443 - val_precision_5: 0.7627 - val_recall_5: 0.9574
Epoch 1/5
30/30 - 8s - 282ms/step - accuracy: 0.4890 - loss: 0.8252 - precision_6: 0.4956 - recall_6: 0.4638 - val_accuracy: 0.5377 - val_loss: 0.6876 - val_precision_6: 0.4894 - val_recall_6: 0.9787
Epoch 2/5
30/30 - 5s - 171ms/step - accuracy: 0.5142 - loss: 0.7098 - precision_6: 0.5200 - recall_6: 0.5383 - val_accuracy: 0.4434 - val_loss: 0.6918 - val_precision_6: 0.4434 - val_recall_6: 1.0000
Epoch 3/5
30/30 - 4s - 121ms/step - accuracy: 0.5551 - loss: 0.6982 - precision_6: 0.5669 - recall_6: 0.5176 - val_accuracy: 0.4434 - val_loss: 0.6913 - val_precision_6: 0.4434 - val_recall_6: 1.0000
Epoch 4/5
30/30 - 4s - 150ms/step - accuracy: 0.5509 - loss: 0.7018 - precision_6: 0.5549 - recall_6: 0.5756 - val_accuracy: 0.5566 - val_loss: 0.6796 - val_precision_6: 0.0000e+00 - val_recall_6: 0.0000e+00
Epoch 5/5
30/30 - 7s - 222ms/step - accuracy: 0.5677 - loss: 0.6890 - precision_6: 0.5831 - recall_6: 0.5155 - val_accuracy: 0.6132 - val_loss: 0.6794 - val_precision_6: 0.5357 - val_recall_6: 0.9574
```

```
In [16]: deep_results
```

```
Out[16]: [ModelResult(name='LSTM', accuracy=0.498113214969635, precision=0.498113214969635, recall=1.0, f1=0.6649874077171052),
ModelResult(name='LSTM + GloVe', accuracy=0.498113214969635, precision=0.498113214969635, recall=1.0, f1=0.6649874077171052),
ModelResult(name='CNN', accuracy=0.849056601524353, precision=0.773809552192688, recall=0.9848484992980957, f1=0.8666666851354765),
ModelResult(name='Transformer', accuracy=0.6301887035369873, precision=0.573913037776947, recall=1.0, f1=0.7292817587192164)]
```

Results

```
In [17]: all_results = ml_results + deep_results
```

```
In [18]: print("\nEvaluation Results:")
for result in all_results:
    print(f"{result.name}: Accuracy={result.accuracy:.4f}, "
          f"Precision={result.precision:.4f}, Recall={result.recall:.4f}, "
          f"F1-Score={result.f1:.4f}")
```


Evaluation Results:

Logistic Regression: Accuracy=0.9547, Precision=0.9615, Recall=0.9470, F1-Score=0.9542

Linear SVM: Accuracy=0.9547, Precision=0.9615, Recall=0.9470, F1-Score=0.9542

Random Forest: Accuracy=0.9321, Precision=0.9524, Recall=0.9091, F1-Score=0.9302

LSTM: Accuracy=0.4981, Precision=0.4981, Recall=1.0000, F1-Score=0.6650

LSTM + GloVe: Accuracy=0.4981, Precision=0.4981, Recall=1.0000, F1-Score=0.6650

CNN: Accuracy=0.8491, Precision=0.7738, Recall=0.9848, F1-Score=0.8667

Transformer: Accuracy=0.6302, Precision=0.5739, Recall=1.0000, F1-Score=0.7293

```
In [20]: def plot_results(results: List[ModelResult], title: str = 'Model Performance Comparison') -> None:
    metrics_df = pd.DataFrame([
        {
            'Model': r.name,
            'Accuracy': r.accuracy,
            'Precision': r.precision,
            'Recall': r.recall,
            'F1-Score': r.f1,
        }
        for r in results
    ])
    metrics_df_melted = metrics_df.melt(id_vars='Model', var_name='Metric', value_name='Score')
    plt.figure(figsize=(10, 6))
    sns.barplot(data=metrics_df_melted, x='Model', y='Score', hue='Metric')
    plt.title(title)
    plt.ylim(0, 1)
    plt.legend(loc='lower right')
    plt.tight_layout()
    plt.savefig('model_performance_comparison.png')
    plt.close()
```

```
In [23]: # Plot comparison chart
plot_results(all_results)
```

```
In [24]: from IPython.display import Image
Image(filename='model_performance_comparison.png')
```

Out[24]:

